

Analytic Geometry

DEF (Dot Product) The dot product between two vectors u and v is defined as

$$u^\top v = u_1v_1 + u_2v_2 + \dots + u_nv_n = \sum_{i=1}^n u_iv_i$$

DEF (Bilinearity) The dot product is bilinear, i.e. for any vectors u, v, w and scalar a ,

$$\begin{aligned} au^\top v &= au^\top v = u^\top av \\ u + v^\top w &= u^\top w + v^\top w \\ w^\top u + v &= w^\top u + w^\top v \end{aligned}$$

DEF (Commutativity) The dot product is commutative, i.e. $u^\top v = v^\top u$

DEF (Inner Product) The inner product between two vectors u and v is defined as

$$\langle u, v \rangle$$

Dot product is a special case of inner product.

DEF (ℓ_2 Norm) The ℓ_2 norm of a vector v is defined as

$$\|v\|_2 = \sqrt{v^\top v} = \sqrt{v_1^2 + v_2^2 + \dots + v_n^2}$$

Also called the Euclidean norm.

DEF (ℓ_2 properties) For all vectors u, v and scalar a ,

- The ℓ_2 norm is non-negative, i.e. $\|v\|_2 \geq 0$.
- $\|au\|_2 = |a| \|u\|_2$ for any scalar a .
- $\|u\|_2$ is zero if and only if u is the zero vector.
- The triangle inequality holds, i.e. $\|u + v\|_2 \leq \|u\|_2 + \|v\|_2$.
- $\|x - y\|_2 = \|y - x\|_2$, also called symmetry.
- $\|u + v\|_2^2 = \|u\|_2^2 + 2u^\top v + \|v\|_2^2$
- $\cos \theta = \frac{u^\top v}{\|u\|_2 \|v\|_2}$ (can be proved using the law of cosines)

THM (Cauchy-Schwarz Inequality) For any vectors u, v , the following inequality holds:

$$|\langle u, v \rangle| \leq \|u\|_2 \|v\|_2$$

DEF (Line) A line is a set of points

$$\{x : x = u + tv \text{ for some } t \in \mathbb{R}\}$$

where u is a point on the line and $v \neq 0$ is the direction vector.

DEF (Plane) A plane is a set of points

$$\{x : v^\top x - u = 0\}$$

where v is the normal vector to the plane and u is the shift from the origin.

DEF (Angle between two vectors) The angle θ between two vectors u and v is given by

$$\cos \theta = \frac{u^\top v}{\|u\|_2 \|v\|_2}$$

DEF (Projection) The vector $\|u\|_2 \cos \theta \frac{v}{\|v\|_2}$ is a projection of u onto v .

DEF (Distance between a point and a plane) The distance between a point z and a plane $v^\top x - u = 0$ is given by

$$\frac{|v^\top z - u|}{\|v\|_2}$$

DEF (Distance between a point and a line) The distance between a point z and a line $x = u + tv$ is given by

$$\left\| z - u - \frac{z - u^\top v}{\|v\|_2^2} v \right\|_2$$

DEF (Singular Value Decomposition (SVD)) The SVD of a matrix X is $U\Sigma V^\top$, where $U^\top U = I$, $V^\top V = I$ and Σ is a diagonal matrix:

$$\Sigma = \begin{bmatrix} \sigma_1 & 0 & \dots & 0 \\ 0 & \sigma_2 & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & \sigma_n \end{bmatrix}$$

and $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_n \geq 0$.

THM (Eckart-Young Theorem) Let $\Sigma_k = \text{diag}(\sigma_1, \dots, \sigma_k, 0, \dots, 0)$ where $k \leq d$. The matrix $U\Sigma_k V^\top$ is the optimal solution to the following problem:

$$\min_{\hat{X}} \|X - \hat{X}\|_F \text{ s.t. } \text{rank}(\hat{X}) \leq k$$

The matrices $Z = U\Sigma_k$ and $W = V^\top$ are the optimal solution to

$$\min_{Z, W} \|X - ZW\|_F \text{ s.t. } Z \in \mathbb{R}^{n \times k}, W \in \mathbb{R}^{k \times d}$$

THM (Useful Identities)

- $A(A^\top A)^{-1} = (AA^\top)^{-1} A^\top$

Calculus

DEF (Derivative) The derivative of a function $f : \mathbb{R} \rightarrow \mathbb{R}$ at x_0 is defined as:

$$(D_x f)(x_0) = \left(\frac{d}{dx} f \right) (x_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + h) - f(x_0)}{h}$$

DEF (Directional Derivative) The directional derivative of a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ along the direction v at $x_0 \in \mathbb{R}^d$ is defined as:

$$(D_v f)(x_0) = \lim_{h \rightarrow 0} \frac{f(x_0 + hv) - f(x_0)}{h}$$

DEF (Partial Derivative) The partial derivative is a directional derivative along the direction of coordinate axes. For a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$, the partial derivative with respect to the i -th coordinate is denoted as:

$$\frac{\partial f}{\partial x_i} = \lim_{h \rightarrow 0} \frac{f(x_1, \dots, x_i + h, \dots, x_d) - f(x_1, \dots, x_i, \dots, x_d)}{h}$$

DEF (Gradient) The gradient of a function is the vector consisting of all partial derivatives denoted by ∇f or $\frac{\partial f}{\partial x}$. For a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$,

EX (Gradient Identities)

- Given a function $f(a) = b^\top a$, $(\nabla f)(a) = b$
- Given a function $f(a) = \|a\|_2^2$, $(\nabla f)(a) = 2a$
- Given a function $f(a) = a^\top Aa$, $(\nabla f)(a) = (A + A^\top)a$

DEF (Hessian) The Hessian of a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is the matrix of second partial derivatives:

$$\nabla^2 f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 x_2} & \dots & \frac{\partial^2 f}{\partial x_1 x_d} \\ \frac{\partial^2 f}{\partial x_2 x_1} & \frac{\partial^2 f}{\partial x_2^2} & \dots & \frac{\partial^2 f}{\partial x_2 x_d} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_d x_1} & \frac{\partial^2 f}{\partial x_d x_2} & \dots & \frac{\partial^2 f}{\partial x_d^2} \end{bmatrix}$$

EX (Hessian Example) Given a function $f(x, y) = x^2 - y^2$, the Hessian is:

$$\nabla^2 f = \begin{bmatrix} 2 & 0 \\ 0 & -2 \end{bmatrix}$$

The 2 indicates that it looks like a cup along the x -axis and the -2 indicates that it looks like an upside-down cup along the y -axis.

Optimisation

DEF (Minimum) The minimum of a function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is written as $\min_x f(x)$, and has the property that $\min_x f(x) \leq f(y)$ for all $y \in \mathbb{R}^d$. The value x^* such that $f(x^*) = \min_x f(x)$ is called the minimizer.

DEF (Convexity) A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is convex if for any $0 \leq \alpha \leq 1$, we have

$$f(\alpha x + (1 - \alpha)y) \leq \alpha f(x) + (1 - \alpha)f(y)$$

In general,

$$f\left(\sum_{i=1}^k \alpha_i x_i\right) \leq \sum_{i=1}^k \alpha_i f(x_i)$$

DEF (Concavity) A function f is concave if $-f$ is convex.

THM (First Order Condition (Convexity)) If f is convex then,

$$\bullet f(x) \geq f(y) + \nabla f(y)^\top (x - y)$$

DEF (Positive Semidefinite) A matrix A is positive semidefinite if for all vectors $v \neq 0$ we have $v^\top A v \geq 0$. Also written as $A \succeq 0$.

THM (Convex Function Implies Positive Semidefinite Hessian) If f is convex, then $\nabla^2 f(x) \succeq 0$.

DEF (Positive Definite) A matrix A is positive definite if for all vectors $v \neq 0$ we have $v^\top A v > 0$. Also written as $A \succ 0$.

DEF (Affine) A function f is affine if $f(x) = Ax + b$ for some matrix A and vector b .

THM (Affine Transform Preservation) If f is convex, then $g(x) = f(Ax + b)$ is also convex.

THM (Non-negative Weighted Sum) If f_1, f_2, \dots, f_k are convex functions, then $f(x) = \sum_{i=1}^k \beta_i f_i(x)$ is convex for all $\beta_i \geq 0$.

EX (Gradient of Quadratic Form) $\nabla_x (x^\top A x) = (A + A^\top)x$

DEF (Strictly Convex) A function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is called strictly if for $0 \leq \alpha < 1$ we have

$$f(\alpha x + (1 - \alpha)y) < \alpha f(x) + (1 - \alpha)f(y)$$

for any $x \neq y$.

DEF (First Order Condition (Strict Convexity)) If f is strictly convex then,

$$\bullet f(x) > f(y) + \nabla f(y)^\top (x - y)$$

THM (Unique Minimizer) If f is strictly convex, then f has a unique minimizer.

THM (Jensen's Inequality) If a function f is convex,

$$f(\mathbb{E}_{x \sim p}[x]) \leq \mathbb{E}_{x \sim p}[f(x)]$$

DEF (Subgradient) A subgradient at x is a vector g that satisfies

$$f(y) \geq f(x) + g^\top (y - x)$$

for any y , and the set of subgradients at x is denoted as $\partial f(x)$. $\nabla f(x) \in \partial f(x)$ if f is differentiable at x . In other words subgradients are tangents that are below the function.

EX (Constrained Optimisation Problem) An example of a constrained optimisation problem is

$$\min_x x^2 \text{ s.t. } -2.5 \leq x \leq -0.5$$

DEF (Feasible Solution) A feasible solution is a point that satisfies all the constraints.

DEF (Lagrangian) If you have an optimisation problem of the form

$$\min_x f(x) \text{ s.t. } h(x) \leq 0$$

the **Lagrangian** is defined as

$$f(x) + \lambda h(x)$$

for some $\lambda \geq 0$ (Lagrange multiplier).

ALG (Solving the Lagrangian)

- Solve $g(\lambda) = \min_x [f(x) + \lambda h(x)]$
- Find $\hat{\lambda}$ such that $\min_x [f(x) + \hat{\lambda} h(x)]$ gives a feasible solution
- Suppose \hat{x} is the solution to the above, and $x^* = \arg \min_{x: h(x) \leq 0} f(x)$ (the optimal solution), then

$$f(\hat{x}) = f(\hat{x}) + \hat{\lambda} h(\hat{x}) \leq f(x^*) + \hat{\lambda} h(x^*) \leq f(x^*)$$

Probability

DEF (Gaussian Distribution) We write $x \sim \mathcal{N}(\mu, \Sigma)$ to denote that x is a random variable with mean μ and covariance Σ . It means that the probability density function of x is given by

$$p(x) = \frac{1}{2\pi^{d/2} |\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu)\right)$$

DEF (Statistical Independence) Two variables x and y are independent if

$$p(x, y) = p(x)p(y)$$

Equivalently,

$$p(x|y) = p(x)$$

. The independence of x and y is denoted by $x \perp y$.

DEF (Statistical Independence [general]) If $\{x_1, \dots, x_n\} \perp \{y_1, \dots, y_m\}$ then

$$p(x_1, \dots, x_n, y_1, \dots, y_m) = p(x_1, \dots, x_n)p(y_1, \dots, y_m)$$

EX (Factorisation of a joint distribution) Suppose $x \in \mathcal{X}, y \in \mathcal{Y}, z \in \mathcal{Z}$. If $\{x, y\} \perp z$ then

$$p(x, y, z) = p(x, y)p(z)$$

The original joint distribution (of size $|\mathcal{X}| \times |\mathcal{Y}| \times |\mathcal{Z}|$) can be factorised into two distributions of size $|\mathcal{X}| \times |\mathcal{Y}|$ and $|\mathcal{Z}|$.

DEF (Mutual Independence) A set of variables $\{x_1, \dots, x_n\}$ are mutually independent if

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i)$$

DEF (Pairwise Independence) A set of variables $\{x_1, \dots, x_n\}$ are pairwise independent if

$$p(x_i, x_j) = p(x_i)p(x_j)$$

for all $i \neq j$.

THM (Mutual Independence implies Pairwise Independence) If a set of variables $\{x_1, \dots, x_n\}$ are mutually independent, then they are pairwise independent. Converse is not true!

DEF (Conditional Independence) The variables x and y are conditionally independent given z if

$$p(x, y|z) = p(x|z)p(y|z)$$

This is denoted by $x \perp y|z$.

DEF (Marginalisation) The marginal distribution of x is obtained by summing out all other variables.

$$p(x) = \sum_y p(x, y)$$

$$p(x|z) = \sum_y p(x, y|z)$$

$$p(x, y|z) = p(x|z)p(y|z)$$

DEF (Bayes Rule) Bayes rule is a way to invert conditional probabilities.

$$p(x|y) = \frac{p(y|x)p(x)}{p(y)}$$

DEF (Chain Rule) Any joint distribution can be factorised into a product of conditional distributions.


$$p(x_1, \dots, x_n) = p(x_1)p(x_2|x_1)p(x_3|x_1, x_2) \cdots p(x_n|x_1, \dots, x_{n-1})$$

DEF ((directed) Graph Representation) A directed graph is a set of nodes connected by edges. Each vertex is a random variable and each edge represents a direct dependency. It is directed and acyclic (DAG). A distribution factorises according to the graph if


$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | \text{parents}(x_i))$$

DEF (Graph structures)


- Chain $x \perp z|y$



- Common cause $x \perp z|y$



- v-structure $x \perp z$ but $x \not\perp z|y$



Training

DEF (Loss Function) Given a predicted output \hat{y} and observed output y , the loss function measures how close the model's prediction is

DEF (Zero-One Loss) $L(y, \hat{y}) = \mathbb{I}_{\{y \neq \hat{y}\}}$

DEF (Mean Squared Error (MSE)) $L(y, \hat{y}) = (y - \hat{y})^2$

DEF (Hinge Loss) $L(y, \hat{y}) = \max(0, 1 - y \cdot \hat{y})$

DEF (Gradient Descent) The gradient descent algorithm is given by

$$x_{t+1} = x_t - \eta \nabla f(x_t)$$

where η is the learning rate/step size.

DEF (Convergence) Given $\epsilon > 0$, we say that an algorithm converges to a point x^* if

$$f(x_t) - f(x^*) \leq \epsilon$$

DEF (Convergence Rate) The convergence rate of an algorithm is the rate at which the algorithm converges to the optimal point. There are three types of convergence rates:

- Sublinear: $f(x_t) - f(x^*) \leq \frac{c}{t^2}$ ($\epsilon = O(\frac{1}{t^2}), t = O(\frac{1}{\sqrt{\epsilon}})$)

- Linear: $f(x_t) - f(x^*) \leq cr^t$ ($\epsilon = O(r^t), t = O(\log \frac{1}{\epsilon})$)
- Quadratic: $f(x_t) - f(x^*) \leq cr^{2^t}$

ALG (Stochastic Gradient Descent (SGD)) Sample a random point x_t, y_t from the dataset and compute the gradient at that point. Repeat until solution is satisfactory.

ALG (Mini-batch Gradient Descent) Sample a mini-batch of points x_t, y_t from the dataset and compute the gradient at that point. Repeat until solution is satisfactory.

DEF (Training) The act of minimising the loss function by adjusting the model's parameters is known as training.

Regression

DEF (Regression) The learning of relationships between input variables x and a numerical output y

DEF (Feature Transformation) We call ϕ a feature transformation. It transforms the input space X into a new space Z .

EX (Feature Transformation Examples)

- Degree 2: $\phi(x) : [1, x_1, x_2] \rightarrow [1, x_1, x_2, x_1^2, x_2^2, x_1x_2]$ (quadratic)

DEF (Linear Regression) Given a dataset $S = \{(\phi(x_1), y_1), \dots, (\phi(x_n), y_n)\}$, minimise the MSE loss function

$$L = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

where $\hat{y}_i = w^T \phi(x_i)$

DEF (Closed-form Solution) The closed-form solution to linear regression is given by

$$w = (\Phi \Phi^T)^{-1} \Phi y$$

DEF (Probabilistic Interpretation) The probabilistic interpretation of linear regression is that

$$y = w^T \phi(x) + \epsilon_i$$

where $\epsilon_i \sim \mathcal{N}(0, 1)$ which implies $y_i \sim \mathcal{N}(w^T \phi(x_i), 1)$. So the log likelihood (L) of the data is

$$L = \sum_{i=1}^N \left[-\frac{1}{2} \log(2\pi) - \frac{1}{2} (y_i - w^T \phi(x_i))^2 \right]$$

Classification

DEF (Classification) Given a dataset S , the goal is to learn a function h that maps input variables x to y , where y is a class label.

DEF (Linear Separability) A dataset is linearly separable if there exists a hyperplane that can separate the data points

DEF (Training of Classification) Find parameters θ such that the zero-one loss is minimised

DEF (Logistic Regression (binary)) A linear classifier that models the probability of a binary class label. The model is given by

$$p(y|x, \theta) = \sigma(-y\theta^T x) = \frac{1}{1 + \exp(-y(w^T x + b))}$$

where σ is the sigmoid function. The classification is done by

$$h(x) = \text{sgn}(w^T x + b)$$

DEF (Maximum Likelihood Estimation (MLE)) The MLE of the parameters θ is given by

$$\hat{\theta} = \arg \max_{\theta} \prod_{i=1}^n p(y_i | x_i, \theta)$$

DEF (Log Likelihood) The log likelihood is applying log to the MLE to obtain a more numerically stable solution:

$$L = \sum_{i=1}^n \log p(y_i | x_i, \theta)$$

ALG (*Training of Logistic Regression*) The training of logistic regression is done by maximising the log likelihood L of w and b .

$$L = \sum_{i=1}^n \log p(y_i | x_i, \theta)$$

There are no closed-form solutions to this problem, so iterative methods like gradient ascent are used which is equivalent to minimising the negative log likelihood.

DEF (*Multiclass Classification*) Using the softmax function, we can extend logistic regression to multiclass classification. Softmax is defined as

$$\text{softmax} \left(\begin{bmatrix} z_1 \\ z_2 \\ \vdots \\ z_k \end{bmatrix} \right) = \begin{bmatrix} \frac{e^{z_1}}{\sum_{i=1}^k e^{z_i}} \\ \frac{e^{z_2}}{\sum_{i=1}^k e^{z_i}} \\ \vdots \\ \frac{e^{z_k}}{\sum_{i=1}^k e^{z_i}} \end{bmatrix}$$

DEF (*Logistic Regression (multiclass)*) The model is given by

$$p(y|x, \theta) = \frac{\exp(\theta_y^T x)}{\sum_{i=1}^k \exp(\theta_i^T x)}$$

The classification is done by

$$h(x) = \underset{y}{\operatorname{argmax}} \theta_y^T x$$

DEF (*Support Vector Machine (SVM) (hard margin)*) A linear classifier that finds the hyperplane that maximises the margin between the classes. It does so by solving the following optimisation problem:

$$\min_w \frac{1}{2} \|w\|_2^2 \quad \text{s.t.} \quad y_i(w^T x_i + b) \geq 1$$

So the lagrangian is

$$L(\alpha, w) = \frac{1}{2} \|w\|_2^2 - \sum_{i=1}^n \alpha_i (y_i (w^T x_i + b) - 1)$$

Solving the lagrangian yields

$$w = \sum_{i=1}^n \alpha_i y_i \cdot \underbrace{x_i}_{\text{support vectors}}$$

The classification is done by

$$h(x) = \operatorname{sgn}(w^T x + b) = \operatorname{sgn} \left(\sum_{i=1}^n \alpha_i y_i x_i^T x + b \right)$$

DEF (*SVM (soft margin)*) Add tolerance to the margin to allow for misclassification. The optimisation problem becomes

$$\min_{w, b} (w^\top w + C \sum_{i=1}^n \xi_i) \quad \text{s.t.} \quad y_i (w^\top x_i + b) \geq 1 - \xi_i$$

The constraint is also known as the hinge loss so when solving the lagrangian

DEF (*Kernel*) A kernel $k : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$ is defined as

$$k(x, x') = \phi(x)^T \phi(x')$$

for some feature function ϕ . It is a measure of similarity between two points in the input space.

EX (*Common Kernels*)

- Linear: $k(x, x') = x^T x'$
- Polynomial: $k(x, x') = (x^T x' + 1)^d$
- Gaussian: $k(x, x') = \exp \left(-\frac{\|x - x'\|_2^2}{2\sigma^2} \right)$

- Hyperbolic Tangent: $k(x, x') = \tanh(kx^T x' + c)$

DEF (*Making Kernels*) A kernel should satisfy

- $k(x, z) = \langle \phi(x), \phi(z) \rangle = \langle \phi(z), \phi(x) \rangle = k(z, x)$ (symmetric)
- $k(z, x)^2 = k(x, x)k(z, z)$
-

$$K = \begin{bmatrix} k(x_1, x_1) & \cdots & k(x_1, x_n) \\ \vdots & \ddots & \vdots \\ k(x_n, x_1) & \cdots & k(x_n, x_n) \end{bmatrix}$$

is positive semi-definite

THM (*Mercer's Theorem*) Suppose k is a continuous symmetric non-negative definite kernel, then k can be expressed as

$$k(x, z) = \sum_{i=1}^{\infty} \lambda_i \phi_i(x) \phi_i(z)$$

where $\{\phi\}$ are eigen-functions, $\|\phi_i\|_2 = 1$ and $\{\lambda_i\}$ are positive eigen values $\lambda_i \geq 0$.

DEF (*Kernel Trick*) The kernel trick is a method to transform the input space into a higher dimensional space without explicitly computing the transformation. This is done by using a kernel function

$$k(x, x') = \phi(x)^T \phi(x')$$

where ϕ is the transformation function.

Neural Networks

DEF (*Perceptron (Single-layer Neural Network)*) A simple linear classifier that learns a weight vector w to classify data points. The perceptron training algorithm can be used to learn the weights. Can model logical functions like AND, OR and NOT.

THM (*Universal Approximation Theorem*) A single-output node NN with a single hidden layer with finite neurons can app finite neurons can approximate continuous (and discontinuous) functions

DEF (*Multi-layer Perceptron (MLP) / Feedforward NN*) A neural network with multiple layers of perceptrons. The network is feedforward as there are no cycles in the network. It can model complex decision boundaries (piecewise linear) [XOR, etc] The perceptron training algorithm can not be used to learn the weights.

The notation $w_{ij}^{(l)}$ denotes the weight from the i th neuron in layer $l - 1$ to the j th neuron in layer l .

DEF (*Activation Function*) Allows for non-linear decision boundaries. They should be differentiable. Also controls the output range to a specific range. Common activation functions are:

- Sigmoid: $\sigma(x) = \frac{1}{1 + e^{-x}}$
- Hyperbolic Tangent (tanh): $\tanh(x) = \frac{1 - \exp(-2x)}{1 + \exp(-2x)}$
- ReLU: $f(x) = \max(0, x)$ (faster than tanh)
- Leaky ReLU solves the 'dying ReLU' problem

DEF (*Computation Graph*) Represents computation as a directed graph comprising of simple operations on vectors and matrices which allows for automatic differentiation.

DEF (*Training of Single-Layer Network*) Given an error function $E_n = \frac{1}{2} (y_n - \hat{y}_n)^2$ (MSE), where $\hat{y}_n = g(a_{nk})$ (some activation function g) [usually sigmoid], and $a_{nk} = w_k^T x_n$ (weighted sum of inputs), the weights are updated using the gradient descent algorithm

$$w_{t+1} = w_t - \eta \frac{\partial E_n}{\partial w_k}$$

By chain rule,

$$\frac{\partial E_n}{\partial w_k} = \frac{\partial E_n}{\partial \hat{y}_n} \frac{\partial \hat{y}_n}{\partial a_{nk}} \frac{\partial a_{nk}}{\partial w_k}$$

0.1. Generalisation

DEF (*Independently and Identically Distributed (i.i.d.)*) A dataset is said to be i.i.d. if the data points come from the same distribution and are statistically independent of each other.

DEF (*Training Error*) For a training set S , the training error for a loss ℓ and a program h is defined as

$$L_S(h) = \frac{1}{n} \sum_{i=1}^n \ell(y_i, h(x_i))$$

DEF (*Generalisation Error*) For a program h , the generalisation error is defined as

$$L_D(h) = \mathbb{E}_{(x,y) \sim \mathcal{D}} [\ell(y, h(x))]$$

The goal of learning is to minimise the generalisation error.

DEF (*Hypothesis Class*) A hypothesis class \mathcal{H} is a set of possible programs of a particular form.

EX (*Hypothesis Class Example*) The set of all linear functions is a hypothesis class i.e.

$$\mathcal{H}_{lin} = \{h(x) = w^T x \mid w \in \mathbb{R}^d\}$$

DEF (*Learning Algorithm*) A learning algorithm is a function that takes a data set of size m and returns a function from the hypothesis class \mathcal{H}

THM (*Probably Approximately Correct (PAC)*) A hypothesis class \mathcal{H} is PAC-learnable with a learning algorithm A if for any distribution \mathcal{D} , and any $\epsilon > 0$ and $0 \leq \delta \leq 1$, there exists $N > 0$ such that for any $n \geq N$:

$$\mathbb{P}_{S \sim \mathcal{D}^n} \left[L_D(A(S)) - \min_{h' \in \mathcal{H}} L_D(h') > \epsilon \right] < \delta$$

In other words, with high probability, the program learned by A achieves a similar error to the best program in \mathcal{H} .

DEF (*Empirical Risk Minimisation (ERM)*) Minimising the loss on a training set is also known as empirical risk minimisation

$$A_{ERM, \mathcal{H}(S)} = h_{ERM} = \underset{h \in \mathcal{H}}{\operatorname{argmin}} L_S(h)$$

THM (*No free lunch theorem*) Suppose $|\mathcal{X}| = 2m$. For any learning algorithm A , there is a distribution \mathcal{D} and $f: \mathcal{X} \rightarrow \{0, 1\}$ such that $L_D(f) = 0$, but

$$\mathbb{P}_{S \sim \mathcal{D}^m} \left[L_D(A(S)) \geq \frac{1}{10} \right] \geq \frac{1}{10}$$

THM (*Error Decomposition*)

$$L_D(h) = \underbrace{L_D(h) - \min_{h' \in \mathcal{H}} L_D(h')}_{\text{Estimation Error}} + \underbrace{\min_{h' \in \mathcal{H}} L_D(h') - L_S(h)}_{\text{Approximation Error}}$$

DEF (*Uniform Convergence*) A hypothesis class \mathcal{H} has uniform convergence property if for any distribution \mathcal{D} , and any $\epsilon > 0$ and $0 \leq \delta \leq 1$, there exists $N > 0$ such that for any $n \geq N$ and every $h \in \mathcal{H}$:

$$\mathbb{P}_{S \sim \mathcal{D}^n} [|L_S(h) - L_D(h)| > \epsilon] < \delta$$

THM (*Uniform Convergence implies PAC learnability*) If a hypothesis class \mathcal{H} has the uniform convergence property, then it is PAC learnable with ERM.

DEF (*Shattering*) A set of n points is shattered by \mathcal{H} if there is an arrangement of n points such that classifiers in \mathcal{H} can produce all 2^n ways of labelling the points.

DEF (*Vapnik-Chervonenkis (VC) Dimension*) VC Dimension is the largest number of points that \mathcal{H} can shatter

DEF (*VC Generalisation Bounds*) With probability $1 - \delta$, for all $h \in \mathcal{H}$:

$$L_D(h) \leq L_S(h) + 2\sqrt{\frac{8d \log(en/d) + 2 \log(4/\delta)}{n}}$$

d is the VC dimension of \mathcal{H} .

EX (*VC Dimension*)

- For linear classifiers, $\text{VC-dim}(\mathcal{H}_{lin}) = p + 1$
- For MLP with p edges, $\text{VC-dim}(\mathcal{H}_{mlp}) = O(p \log p)$

DEF (*Surrogate Loss*) A surrogate loss is a loss function that is easier to optimise than the original loss function. Cross entropy or log likelihood are common surrogate losses for 0-1 loss.

DEF (*Underfitting*) A model h is underfitting if there is another model f that has a lower training i.e. $L_S(f) < L_S(h)$.

DEF (*Overfitting*) A model h is overfitting if there is another model f that has a higher training error (S) but a lower test error (S') i.e. $L_S(f) > L_S(h)$ and $L_{S'}(f) < L_{S'}(h)$.

DEF (*Development Set*) A development set is a subset of the training set that is used to tune hyperparameters.

DEF (*Stability*) A learning algorithm is stable if the learned program does not change much in performance when we change the data set slightly.

THM (*Stability prevents Overfitting*) Stable learning algorithms don't overfit

DEF (λ -strongly convex) A function f is λ -strongly convex if for all x, y and $\lambda > 0$:

$$f(y) \geq f(x) + \nabla f(x)^\top y - x + \frac{\lambda}{2} \|y - x\|_2^2$$

THM (L_2 Regularisation) Regularisation is a technique to prevent overfitting by adding a penalty term to the loss function. For L_2 regularisation, the loss function is modified to

$$L = \frac{1}{n} \sum_{i=1}^n \ell(y_i, h(x_i)) + \frac{\lambda}{2} \|w\|_2^2$$

Note that if the loss function is convex, then the regularised loss function is λ -strongly convex.

DEF (*Hardness of Optimising Neural Networks*) Training a 2-layer 3 node neural network is NP-complete. If we could minimise the loss function in polynomial time, then we would be able to solve the P=NP problem.

DEF (*Overparameterisation*) Overparameterisation is the practice of using more parameters than necessary to fit the training data. It helps with optimisation.

DEF (*Interpolation*) A model interpolates the data if it achieves zero training error.

0.2. PCA

DEF (*Dimensionality Reduction*) Dimensionality reduction is the process of reducing the number of features in a dataset. This is used for visualisation, exploration and compression.

DEF (*Principal Component Analysis (PCA)*) PCA is a method for dimensionality reduction. It finds the directions that maximise the variance in the data. These directions are called principal components. Each principle component is orthogonal to the others.

DEF (*Maximum Variance Formulation*) An optimisation problem that PCA solves is to find the directions that maximise the variance in the data. This can be solved with lagrange multipliers.

Expectation Maximisation (EM)

DEF (*Expectation Maximisation*) If can't solve an optimisation problem directly, EM is used. It yields a lower bound of the original loss. It utilises Jensen's inequality.

DEF (*K-means Distortion*)

$$J = \sum_{n=1}^N \sum_{k=1}^K r_{nk} \|x_n - \mu_k\|^2$$

DEF (*K-means Steps*)

- Randomly initialise $\mu_{k=1, \dots, K}$
- Minimise J with respect to r_{nk} (Expectation)
- Minimise J with respect to μ_k (Maximisation)
- Repeat until convergence

DEF (*K-means Solutions*)

$$\mu_k = \frac{\sum_n r_{nk} x_n}{\underbrace{\sum_n r_{nk}}_{\text{Mean of cluster}}}, (k = \arg \min_j \|x_n - \mu_j\|^2)$$

DEF (*GMM Probabilities*)

$$p(z_k = 1) = \pi_k, p(\mathbf{z}) = \prod_{k=1}^K \pi_k^{z_k}, p(\mathbf{x}|\mathbf{z}) = \prod_{k=1}^K \mathcal{N}(x|\mu_k, \Sigma_k)^{z_k}$$

DEF (*GMM Responsibility Function*)

$$r_{nk} = p(z_k = 1|x_n) = \frac{\pi_k \mathcal{N}(x_n|\mu_k, \Sigma_k)}{\sum_{j=1}^K \pi_j \mathcal{N}(x_n|\mu_j, \Sigma_j)}$$

DEF (*Kullback-Leibler Divergence*) Measure how well one probability distribution approximates another

$$KL(p||q) = \mathbb{E}_{x \sim p} \left[\log \frac{p(x)}{q(x)} \right] = \sum_{x \in X} p(x) \log \frac{p(x)}{q(x)}$$

Not a metric as it is not symmetric $KL(p||q) \neq KL(q||p)$